SDNFV Lab2

Part1

1. How many OpenFlow headers with type “OFPT_FLOW_MOD” and command “OFPFC_ADD” are there among all the packets? What are the match fields and the corresponding actions in each “OFPT_FLOW_MOD” message? What are the Idle Timeout values for all flow rules on s1 in GUI?

   There are 2 distinct “OFPT_FLOW_MOD” headers during the experiment.

Match fields	Actions	Timout values
In packets		
IN_PORT = 1 ETH_DST = da:ad:c8:9a:37:40 ETH_SRC = 1e:30:5f:0e:ab:65	Output port = 2	10
IN_PORT = 2 ETH_DST = 1e:30:5f:0e:ab:65 ETH_SRC = da:ad:c8:9a:37:40	Output port = 1	10
Other rules		
ETH_TYPE = lldp	Output port = CONTROLLER	0
ETH_TYPE = bddp	Output port = CONTROLLER	0
ETH_TYPE = ipv4	Output port = CONTROLLER	0
ETH_TYPE = arp	Output port = CONTROLLER	0

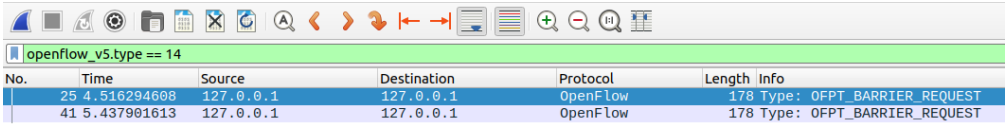


```
Idle timeout: 0
Hard timeout: 0
Priority: 10
Buffer ID: OFP_NO_BUFFER (4294967295)
Out port: OFPP_ANY (4294967295)
Out group: OFPG_ANY (4294967295)
▸ Flags: 0x0001
Importance: 0
▾ Match
    Type: OFPMT_OXM (1)
    Length: 32
  ▾ OXM field
      Class: OFPXMC_OPENFLOW_BASIC (0x8000)
      0000 000. = Field: OFPXMT_OFB_IN_PORT (0)
      .... ...0 = Has mask: False
      Length: 4
      Value: 2
  ▾ OXM field
      Class: OFPXMC_OPENFLOW_BASIC (0x8000)
      0000 011. = Field: OFPXMT_OFB_ETH_DST (3)
      .... ...0 = Has mask: False
      Length: 6
      Value: 1e:30:5f:0e:ab:65 (1e:30:5f:0e:ab:65)
  ▾ OXM field
      Class: OFPXMC_OPENFLOW_BASIC (0x8000)
      0000 100. = Field: OFPXMT_OFB_ETH_SRC (4)
      .... ...0 = Has mask: False
      Length: 6
      Value: da:ad:c8:9a:37:40 (da:ad:c8:9a:37:40)
```

Part 2

```
mininet> h1 arping -c 4 h2
ARPING 10.0.0.2
42 bytes from da:ad:c8:9a:37:40 (10.0.0.2): index=0 time=1.007 msec
42 bytes from da:ad:c8:9a:37:40 (10.0.0.2): index=1 time=25.663 usec
42 bytes from da:ad:c8:9a:37:40 (10.0.0.2): index=2 time=3.838 usec
42 bytes from da:ad:c8:9a:37:40 (10.0.0.2): index=3 time=3.743 usec

--- 10.0.0.2 statistics ---
4 packets transmitted, 4 packets received,   0% unanswered (0 extra)
rtt min/avg/max/std-dev = 0.004/0.260/1.007/0.431 ms
```

```
mininet> h1 ping -c 5 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.039 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.075 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.059 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.049 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.058 ms

--- 10.0.0.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4114ms
rtt min/avg/max/mdev = 0.039/0.056/0.075/0.011 ms
mininet>
```

Part 3

I created the following topology. I created a rule that send all packets with ETH_TYPE = ARP to all the other ports. Once all the switches have this rule and a packet is sent, the swithes will endless sending packets to each other in a cycle, creating a broadcast storm.

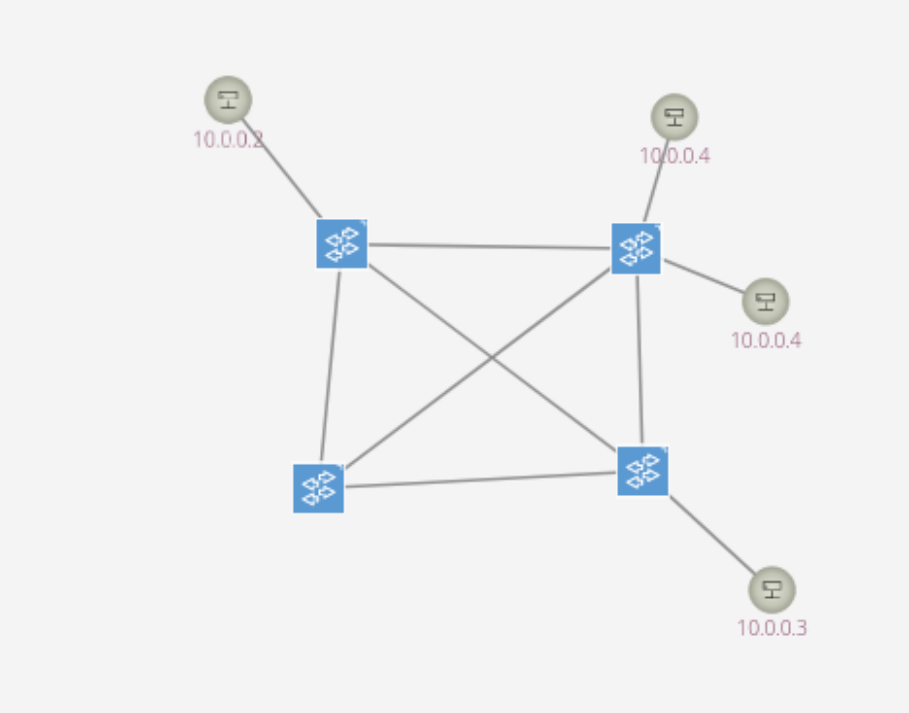

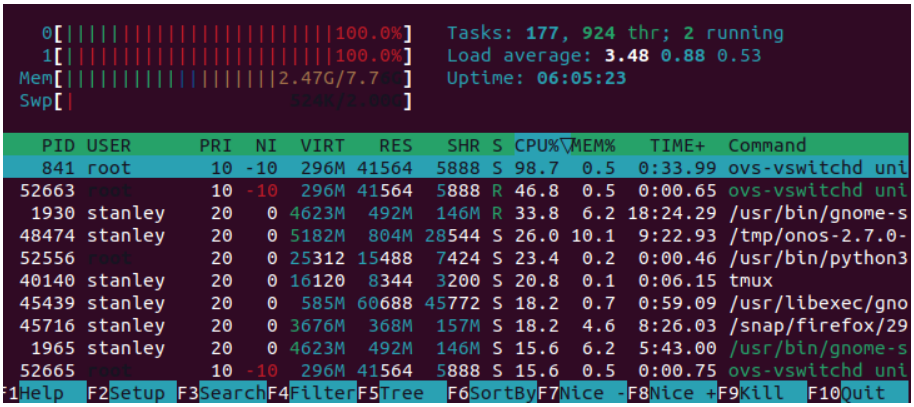

Part 4

1. In data plane, h1 send an ICMP packet to the s1
2. In the control plane, the packet meets the selector ETH_TYPE = ipv4 on the switch, hence forwarded to the controller.
3. The fwd app handles the packet, creating a rule using flowObjectiveService, and forward the packet to the destination host. (There are several steps in the app, such as detecting whether to drop the packet, or checking it is the edge...)

In the source code, the app calls the function to install the flow rule first, then forward packet to the destination. However, wireshark captures the packet received from h2 before capturing the request to add the flow rule. This is because the action of making new flow rules are asychronous, and the packet forwarding is slightly faster than adding new flow rules.

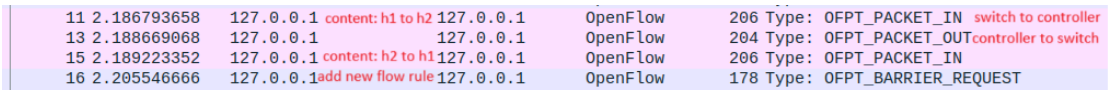

Part 5

In this lab,

1. I learned how to install flow rules with json files using curl and GUI
2. I learned how to capture and inspect packets using wireshark
3. I traced the code the reactive forwarding app.
4. I observed how the controller and switch interacts.